


# Lecture 05

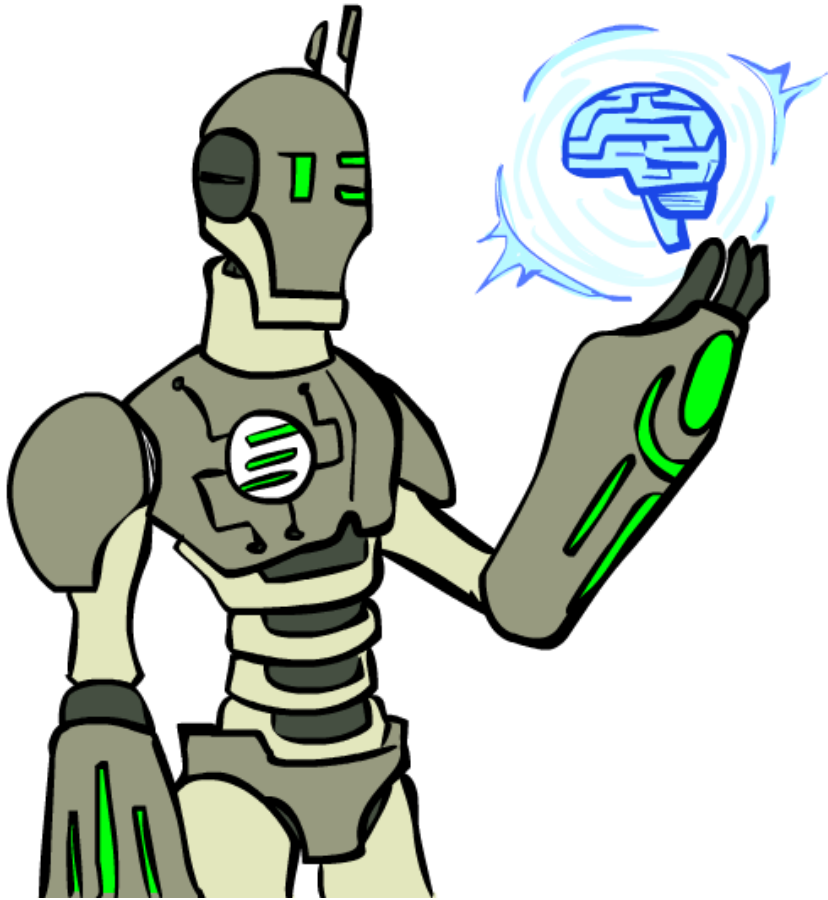
Ashis Kumar Chanda  
[chanda@rowan.edu](mailto:chanda@rowan.edu)



# Today

---

- Minimax algorithm
- Alpha-beta pruning



# Introduction

---

- In previous classes, we learned different searching methods with and without knowledge.
- We had a specific cost for each node in a search tree and goal was fixed.
- Let's consider what happens when we have **opponent**.
- And the opponent is also **rational**.


# Types of Games

---



- General Games

- Agents have independent utilities (values on outcomes)
- Cooperation, indifference, competition, and more are all possible



- Zero-Sum Games

- Agents have opposite utilities (values on outcomes)
- Lets us think of a single value that one maximizes and the other minimizes

# Adversarial Search

---


- In a multi-agent environment, we can take two types of actions:
  - Maximizing utility (without having to predict the action of any individual agent)
  - Explicitly model the adversarial agents with the techniques of adversarial game-tree search.
- "Unpredictable" opponent → specifying a move for every possible opponent's reply.

# Games as Adversarial Search

---

- Let's define a game as a search tree, where
  - Move = Action
  - Board position = State or Node
  - Terminal states = Where the game has ended
  - Utility function = Defines the final numeric value to player  $p$  when the game ends in terminal state
- Special feature:
  - Two players, adversarial

# Two Player Game Example: Tic-tac-toe



Each leaf node indicates the **utility** value of the terminal state from the point of the view of MAX.

Win = 1


Draw = 0

Lose = -1

A game tree for the game of tic-tac-toe. The top node is the initial state, and MAX moves first, placing an X in an empty square.

# Two Player Game Example: Chess

- Suppose, we have two opponents in a chess game, **Red** and **Blue**.
- **Red** moves first.
- We can think about the game moves like the following tree where the leaf nodes are terminal states.




- Red nodes are the winning state for Red.
- Blue nodes are the winning states for blue.
- Gray nodes are draw states.
- Red will try to reach at final red states, and so, we need to think from the **bottom**.




# Two Player Game Example: Chess

---



# Two Player Game Example: Chess


---



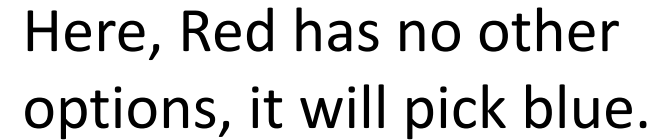
Red will try to reach at final red states

# Two Player Game Example: Chess

---




Red will try to reach at final red states



# Two Player Game Example: Chess


---



Blue will try to reach at final blue states


# Two Player Game Example: Chess

---



Blue will try to reach at final blue states

# Two Player Game Example: Chess




From the initial state,  
Red will select the left  
direction.

# Adversarial Search (Minimax)

- Deterministic, zero-sum games:
  - Tic-tac-toe, chess, checkers
  - One player maximizes result
  - The other minimizes result
- Minimax search:
  - A state-space search tree
  - Players alternate turns
  - Compute each node's **minimax value**: the best achievable utility against a rational (optimal) adversary

**Minimax values:**  
**computed recursively**



**Terminal values:**  
**part of the game**



# Minimax Procedure

---

- If at limit of search, compute static value
  - Relative to player
- If **minimizing** level, do minimax
  - Report minimum
- If **maximizing** level, do minimax
  - Report maximum

# Two Player Games

---

- We have two players, Max and Min.
- Max always moves first.
- Min is the opponent.

We have

- An initial state.
  - A set of actions.
  - A terminal test (which tells us when the game is over).
  - A utility function (evaluation function).
- 
- The utility function is like the heuristic function we have seen in the past, except it evaluates a node in terms of **how good it is for each player.**



Max Vs Min

# Minimax Example

- Idea: choose a move to a position with the highest **minimax value** = best achievable payoff against a **rational** opponent.


Example: deterministic 2-ply game:

Minimax value is computed bottom up:  
-Leaf values are given. **MAX**

-3 is the best outcome for MIN in this branch. **MIN**


-3 is the best outcome for MAX in this game.

-We explore this tree in depth-first manner.



# Minimax Example

---



# Minimax Implementation (Dispatch)

---

```
def value(state):
```

if the state is a terminal state: return the state's utility

if the next agent is MAX: return max-value(state)

if the next agent is MIN: return min-value(state)

```
def max-value(state):
```

initialize  $v = -\infty$

for each successor of state:

$v = \max(v, \text{value}(\text{successor}))$

return  $v$

```
def min-value(state):
```

initialize  $v = +\infty$

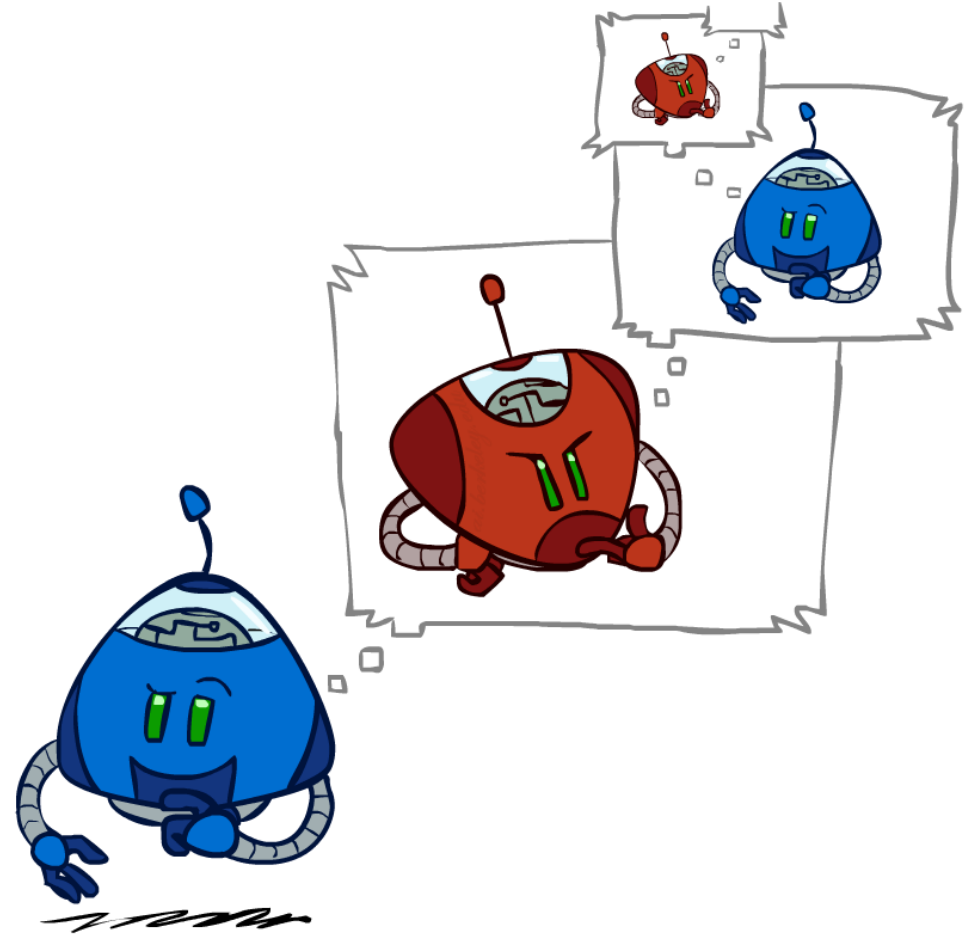
for each successor of state:

$v = \min(v, \text{value}(\text{successor}))$

return  $v$


# Minimax Efficiency

- How efficient is minimax?
  - Just like (exhaustive) DFS
  - Time:  $O(b^m)$
  - Space:  $O(bm)$
- The number of game states is exponential in the depth of the tree.




# Minimax Efficiency

- Example:
  - For tic-tac-toe game, we have  $9! = 362,880$  terminal nodes.
- Example:
  - For chess,  $b \approx 35$ ,  $m \approx 100$
  - Exact solution is completely infeasible
  - But, do we need to explore the whole tree?



# Game Tree Pruning


---







# Simple $\alpha$ - $\beta$ Example


---




# $\alpha$ - $\beta$ pruning example



# $\alpha$ - $\beta$ pruning example




# $\alpha$ - $\beta$ pruning example




MIN will do at least as good as 14 in this branch (which is very good for MAX!) so MAX will want to explore this branch more.

# $\alpha$ - $\beta$ pruning example



MIN will do at least as good as 5 in this branch (which is still good for MAX) so MAX will want to explore this branch more.

# $\alpha$ - $\beta$ pruning example



pity (for MAX): MIN will be able to play this last branch and get 2. This is worse than 3, so MAX will play 3.


# $\alpha$ - $\beta$ pruning

---

- Why we call it as  $\alpha$ - $\beta$  pruning?
- The name come from the two extra parameters in `MaxValue(state,  $\alpha$ ,  $\beta$ )`.
- $\alpha$  = The **highest value** choice we have found so far at choice point along the path for MAX
- $\beta$  = The **lowest value** choice we have found so far at choice point along the path for MIN.

# $\alpha$ - $\beta$ Pruning best and worst case


- Worst case:
  - Bad ordering: Alpha-beta prunes **NO** nodes
- Best case:
  - Assume cooperative oracle orders nodes
    - **Best value on left**






# Alpha-Beta Quiz

---



# Alpha-Beta Quiz 2




# Resource Limits

---



# Resource Limits


- Problem: In realistic games, cannot search to leaves!
- Solution: Depth-limited search
  - Instead, search only to a limited depth in the tree
  - Replace terminal utilities with **an evaluation function** for non-terminal positions
- Example:
  - Suppose we have 100 seconds, can explore 10K nodes / sec
  - So can check 1M nodes per move
  - $\alpha$ - $\beta$  reaches about depth 8 – decent chess program
- Guarantee of optimal play is gone
- More plies makes a BIG difference
- Use iterative deepening for an anytime algorithm



# Next class?

---

- Propositional logic
- First order logic
- Quantifier



Thanks!